

## Homework 1 – [Bhati –xxxxxxxx36]

### Screenshots of Coding Output

My ID = 36. So the last two digits are 3 and 6.

```
# filter out the training and test sets with these two digits - 3 and 6
x_train = x_train[(y_train == 3) | (y_train == 6)]
y_train = y_train[(y_train == 3) | (y_train == 6)]
x_test = x_test[(y_test == 3) | (y_test == 6)]
y_test = y_test[(y_test == 3) | (y_test == 6)]
```

```
# For binary classification modify labels: Label encoding [3 => 0 and 6 => 1]
y_train = np.where(y_train == 3, 0, 1)
y_test = np.where(y_test == 3, 0, 1)
```

```
In [13]: 1 # Pixel values range from 0 to 255. Normalize pixel values of train and test to the range [0, 1]
2 x_train, x_test = x_train / 255.0, x_test / 255.0
```

```
In [14]: 1 x_train[0] # Normalized x_train
0.10980592, 0.          , 0.          , 0.          , 0.          ,
0.          , 0.          , 0.          ], [
0.          , 0.          , 0.          , 0.          , 0.          ,
0.          , 0.          , 0.          , 0.17647059, 0.87058824,
0.98823529, 0.98823529, 0.98823529, 0.98823529, 0.99215686,
0.98823529, 0.98823529, 0.98823529, 0.69411765, 0.          ,
0.          , 0.          , 0.          , 0.          , 0.          ,
0.          , 0.          , 0.          ], [
0.          , 0.          , 0.          , 0.          , 0.          ,
0.          , 0.          , 0.          , 0.17647059, 0.8745098 ,
0.99215686, 0.99215686, 0.99215686, 0.99215686, 1.          ,
0.99215686, 0.99215686, 0.99215686, 0.99215686, 0.29019608,
0.          , 0.          , 0.          , 0.          , 0.          ,
0.          , 0.          , 0.          ], [
0.          , 0.          , 0.          , 0.          , 0.          ,
0.          , 0.          , 0.          , 0.          , 0.12156863,
0.48235294, 0.20392157, 0.17254902, 0.17254902, 0.17254902,
0.17254902, 0.56078431, 0.98823529, 0.98823529, 0.29019608,
0.          , 0.          , 0.          , 0.          , 0.          ,
0.          , 0.          , 0.          ],
```

## Neural Network for Binary Classification

```
[15]: 1 # Batch size and Learning rate are hyperparameters
2 batch_size = 32
3 learning_rate = 0.001

[16]: 1 model = Sequential()
2
3 model.add(Flatten(input_shape=(28,28))) ## 28x28 pixels are converted to 1D 784 input units.
4 model.add(Dense(128,activation='relu')) # 128 nodes with ReLU activation to add non-linearity.
5 model.add(Dense(1,activation='sigmoid')) ## 1 output node for binary classification using sigmoid.

[17]: 1 model.summary()
2

Model: "sequential"
-----  
Layer (type)          Output Shape         Param #
-----  
flatten (Flatten)     (None, 784)          0  
dense (Dense)         (None, 128)          100480  
dense_1 (Dense)       (None, 1)           129  
-----  
Total params: 100,609  
Trainable params: 100,609  
Non-trainable params: 0
```

Params in the second layer are 100480 because 784 inputs \* 128 nodes(weights = 100352) + 128(baises) = 100480

Params in the last layer = 129 because 128(wts) + 1(bias) = 129

**Binary\_Crossentropy** is used for binary classification and optimization is chosen as Adam

```
[18]: 1 # Compiling the model
2 model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=learning_rate), loss='binary_crossentropy',metrics=['accuracy'])
```

**Early stopping criteria based on validation loss**

```
[19]: 1 # Early stopping criteria based on validation loss
2 # patience = Number of epochs with no improvement after which training will be stopped.
3 # restore_best_weights = Whether to restore model weights from the epoch with the best value of the monitored quantity.
4 early_stopping = EarlyStopping(monitor='val_loss', patience=10, restore_best_weights=True)
```

```
[20]: 1 # Training the model
2 history = model.fit(x_train, y_train, batch_size=batch_size, epochs=20, validation_split=0.2, callbacks=[early_stopping])

Epoch 1/20
302/302 [=====] - 1s 2ms/step - loss: 0.0344 - accuracy: 0.9899 - val_loss: 0.0086 - val_accuracy: 0.975
Epoch 2/20
302/302 [=====] - 0s 1ms/step - loss: 0.0098 - accuracy: 0.9969 - val_loss: 0.0055 - val_accuracy: 0.979
Epoch 3/20
302/302 [=====] - 0s 2ms/step - loss: 0.0054 - accuracy: 0.9985 - val_loss: 0.0045 - val_accuracy: 0.979
Epoch 4/20
302/302 [=====] - 0s 2ms/step - loss: 0.0026 - accuracy: 0.9997 - val_loss: 0.0021 - val_accuracy: 0.996
Epoch 5/20
302/302 [=====] - 0s 1ms/step - loss: 0.0011 - accuracy: 0.9999 - val_loss: 0.0082 - val_accuracy: 0.997
Epoch 6/20
```

```

30]: 1 # Evaluate the model on the test data
      2 test_loss, test_accuracy = model.evaluate(x_test, y_test)
      3 print("Test Accuracy:", test_accuracy)

62/62 [=====] - 0s 869us/step - loss: 0.0024 - accuracy: 0.9990
Test Accuracy: 0.9989837408065796

```

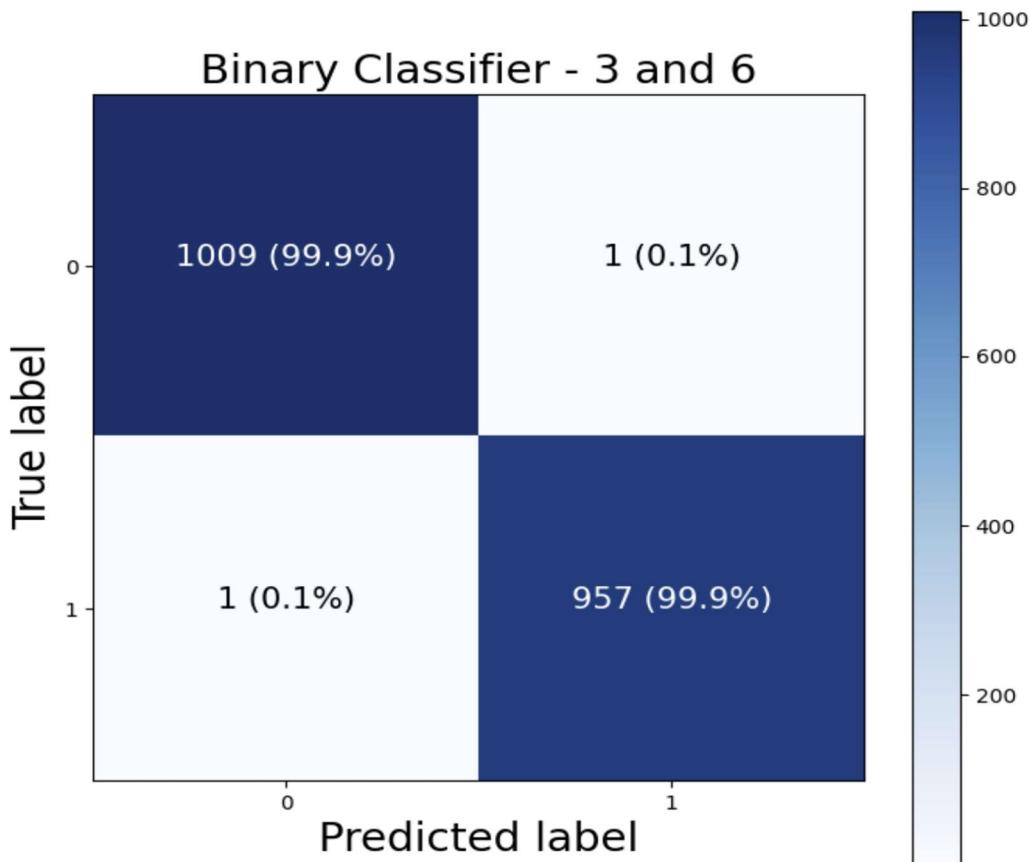
### Construct a confusion matrix

```

28]: 1 conf_matrix = confusion_matrix(y_test, y_pred)
      2 print("Confusion Matrix:")
      3 print(conf_matrix)
      4 draw_confusion_matrix(y_test, y_pred,"Binary Classifier - 3 and 6")

Confusion Matrix:
[[1009    1]
 [    1  957]]

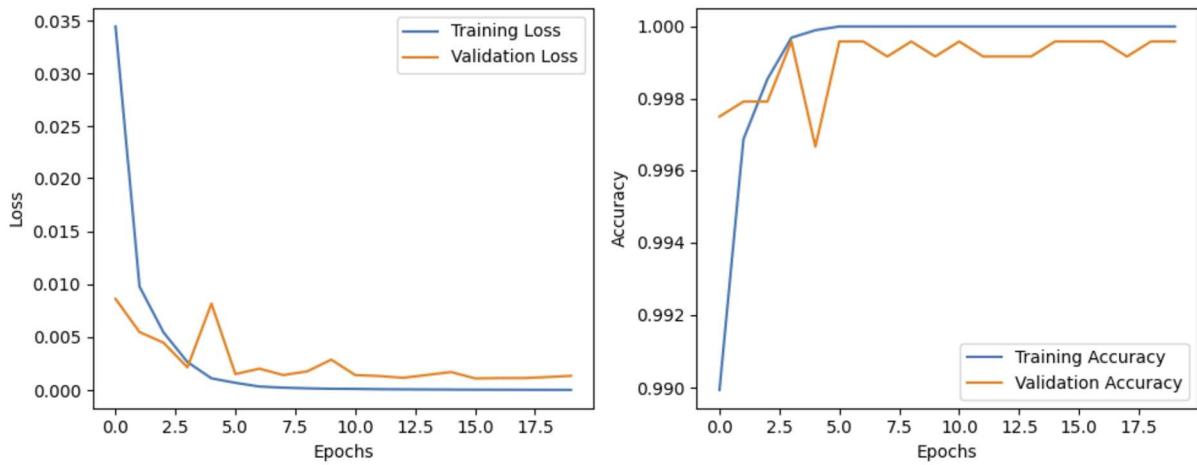
```



Just two instances have been misclassified. 1 instance of class 3 has been misclassified as class 6(False Negative) and 1 instance of Class 6 has been misclassified as 3(False Positive)

### Learning Curve - Loss and Accuracy

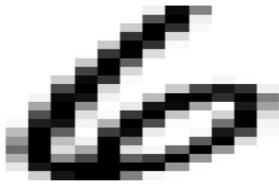
```
.]: 1 # Present the Learning curve
2 plt.figure(figsize=(10, 4))
3 plt.subplot(1, 2, 1)
4 plt.plot(history.history['loss'], label='Training Loss')
5 plt.plot(history.history['val_loss'], label='Validation Loss')
6 plt.xlabel('Epochs')
7 plt.ylabel('Loss')
8 plt.legend()
9
10 plt.subplot(1, 2, 2)
11 plt.plot(history.history['accuracy'], label='Training Accuracy')
12 plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
13 plt.xlabel('Epochs')
14 plt.ylabel('Accuracy')
15 plt.legend()
16
17 plt.tight_layout()
18 plt.show()
```



### Show some examples of predictions

```
: 1 eg_indices = np.random.choice(len(x_test), 10) # 10 randomly selected samples
2 eg_images = x_test[eg_indices]
3 eg_labels = y_test[eg_indices]
4 eg_predictions = model.predict(eg_images)
5 for i in range(len(eg_images)):
6     plt.figure()
7     plt.imshow(eg_images[i], cmap='Greys')
8     plt.title(f"True Label: {eg_labels[i]}, Predicted: {np.round(eg_predictions[i][0])}")
9     plt.axis('off')
10
```

True Label: 1, Predicted: 1.0



True Label: 0, Predicted: 0.0



TABLE for 2-Class(Binary Class) Hyperparameters

Activation function – hidden layer	ReLU
Activation function - output layer	Sigmoid
Weight Initializer	Default -GlorotUniform
Number of hidden layers	1 (2 including output layer)
Neurons in hidden layers	128
Neurons in the output layer	1
Loss function	Binary_crossentropy
Optimizer	Adam
Number of epochs	20
Batch size	32
Learning rate	0.001
Evaluation Metric	Accuracy

B. Three different Weight Initializers. As per keras documentation there are multiple initializers such as He normal, GlorotNormal, RandomNormal, and GlorotUniform.

```
3]: 1 from tensorflow.keras.initializers import GlorotNormal, HeNormal, RandomNormal,GlorotUniform  
      2
```

```

1 # Lists to store results
2 learning_curves = []
3 accuracies = []

1 labels = ["GlorotNormal", "HeNormal", "RandomNormal", "GlorotUniform"]
2 i=0
3 # Build and train models for each initializer
4 for initializer in initializers:
5
6     model = Sequential([
7         Flatten(input_shape=(28, 28)),
8         Dense(128, activation='relu', kernel_initializer=initializer),
9         Dense(1, activation='sigmoid')
10    ])
11    model.summary()
12    model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=learning_rate), loss='binary_crossentropy', metrics=['accuracy'])
13
14    # Early stopping criteria based on validation loss
15    # patience = Number of epochs with no improvement after which training will be stopped.
16    # restore_best_weights = Whether to restore model weights from the epoch with the best value of the monitored quantity.
17    early_stopping = EarlyStopping(monitor='val_loss', patience=10, restore_best_weights=True)
18
19    # Train the model
20    history = model.fit(x_train, y_train, batch_size=batch_size, epochs=20, validation_split=0.2, verbose=0, callbacks=[early_stopping])
21
22    # Evaluate on test data
23    y_pred = np.round(model.predict(x_test))
24    accuracy = accuracy_score(y_test, y_pred)
25    confusion = confusion_matrix(y_test, y_pred)
26    draw_confusion_matrix(y_test, y_pred, labels[i])
27
28    #confusion_matrices.append(confusion)
29    learning_curves.append(history.history)
30    accuracies.append(accuracy)
31    i=i+1

```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
<hr/>		
flatten_1 (Flatten)	(None, 784)	0
dense_2 (Dense)	(None, 128)	100480
dense_3 (Dense)	(None, 1)	129
<hr/>		
Total params: 100,609		
Trainable params: 100,609		
Non-trainable params: 0		

62/62 [=====] - 0s 738us/step

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
<hr/>		
flatten_2 (Flatten)	(None, 784)	0
dense_4 (Dense)	(None, 128)	100480
dense_5 (Dense)	(None, 1)	129
<hr/>		
Total params: 100,609		
Trainable params: 100,609		
Non-trainable params: 0		

62/62 [=====] - 0s 721us/step

Model: "sequential\_3"

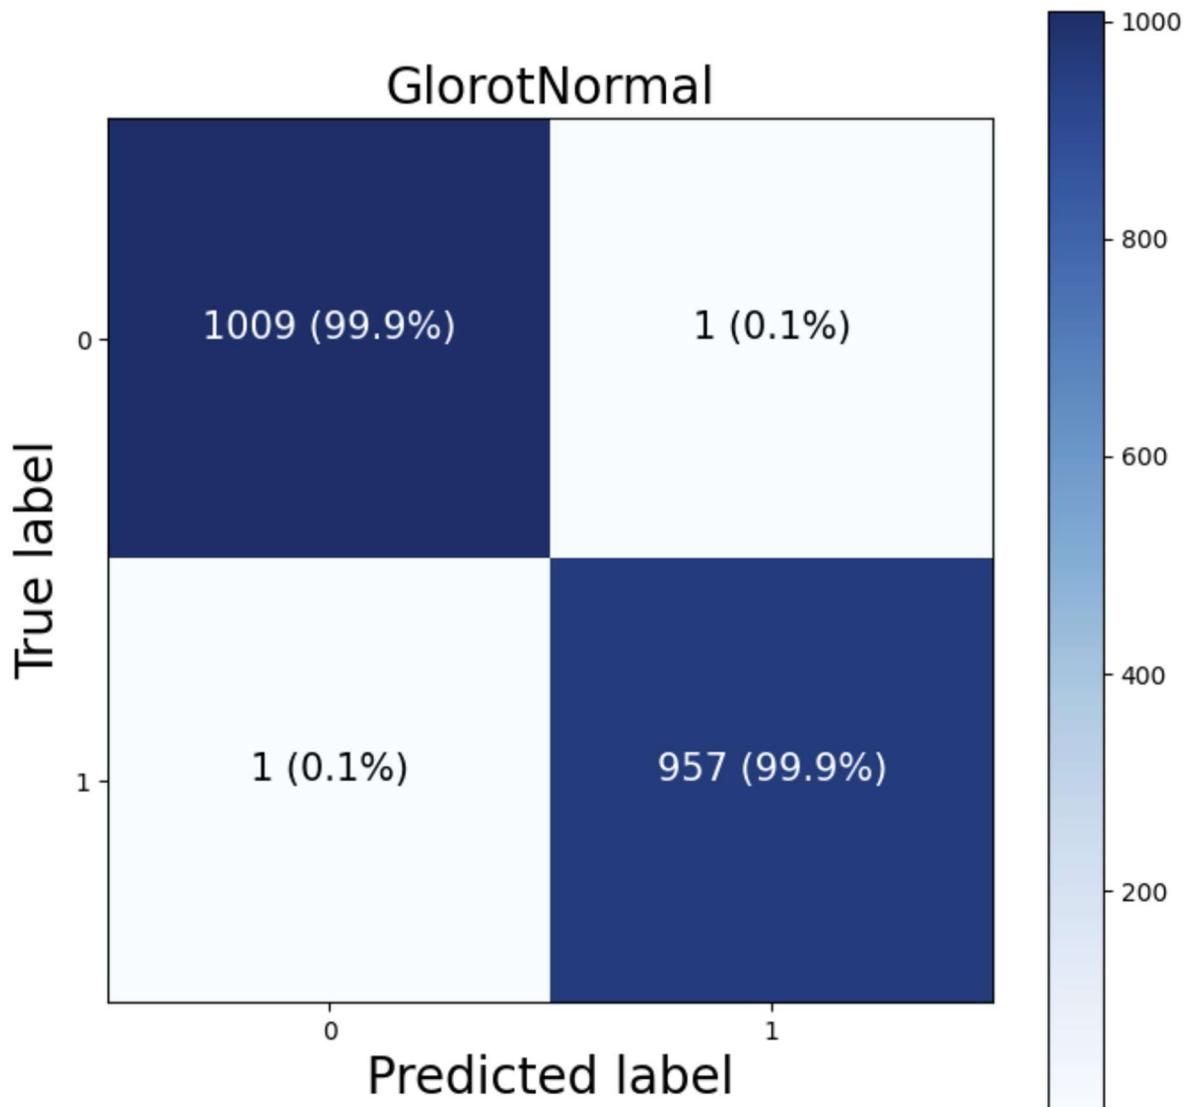
Layer (type)	Output Shape	Param #
<hr/>		
flatten_3 (Flatten)	(None, 784)	0
dense_6 (Dense)	(None, 128)	100480
dense_7 (Dense)	(None, 1)	129
<hr/>		
Total params: 100,609		
Trainable params: 100,609		
Non-trainable params: 0		

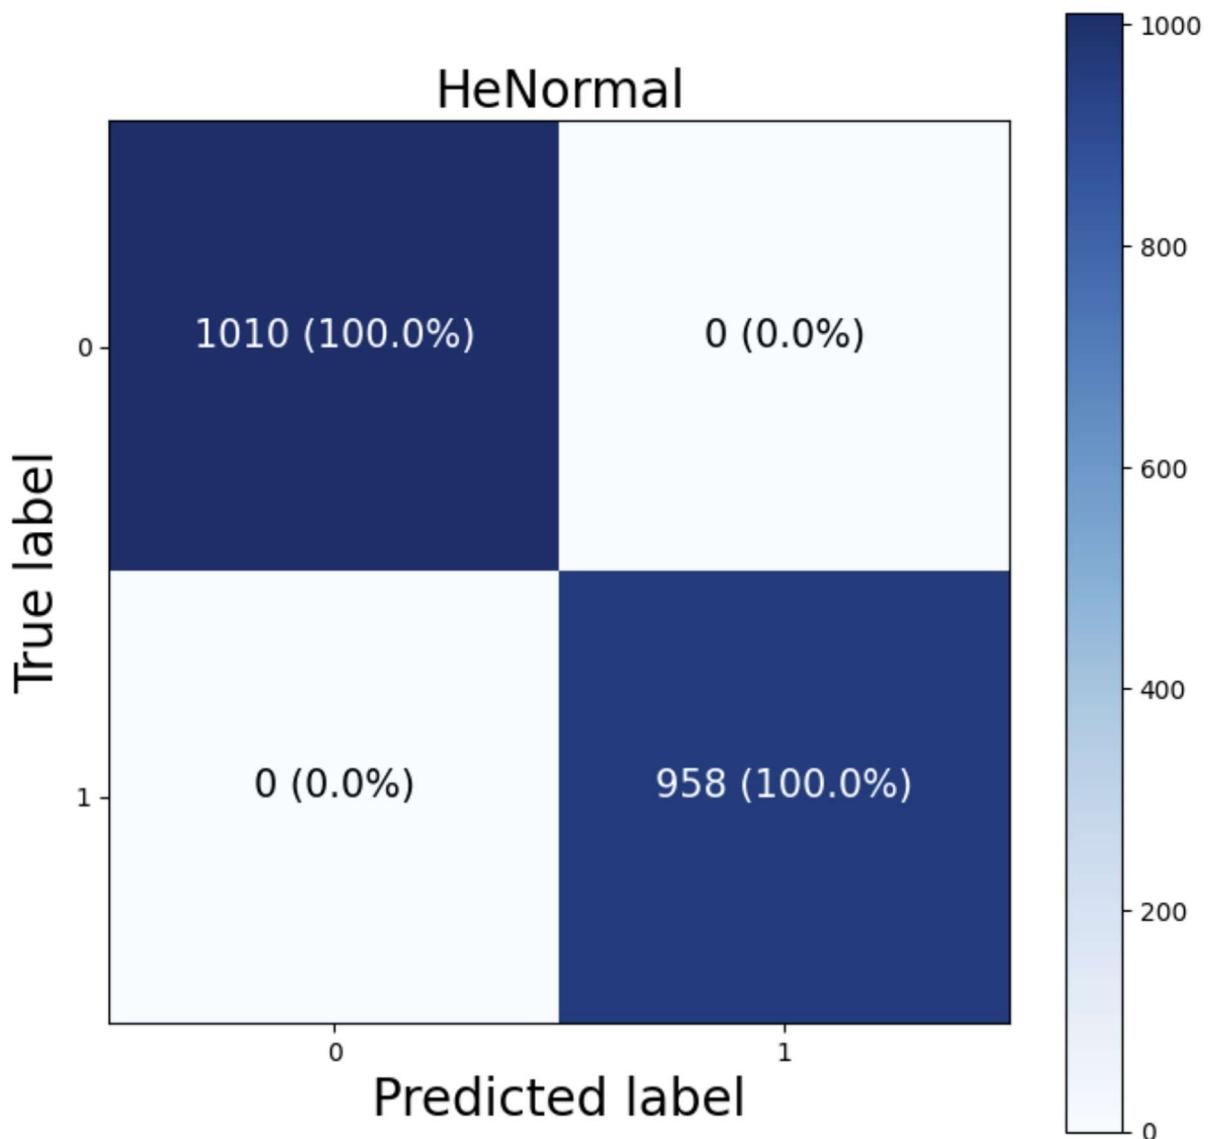
62/62 [=====] - 0s 754us/step

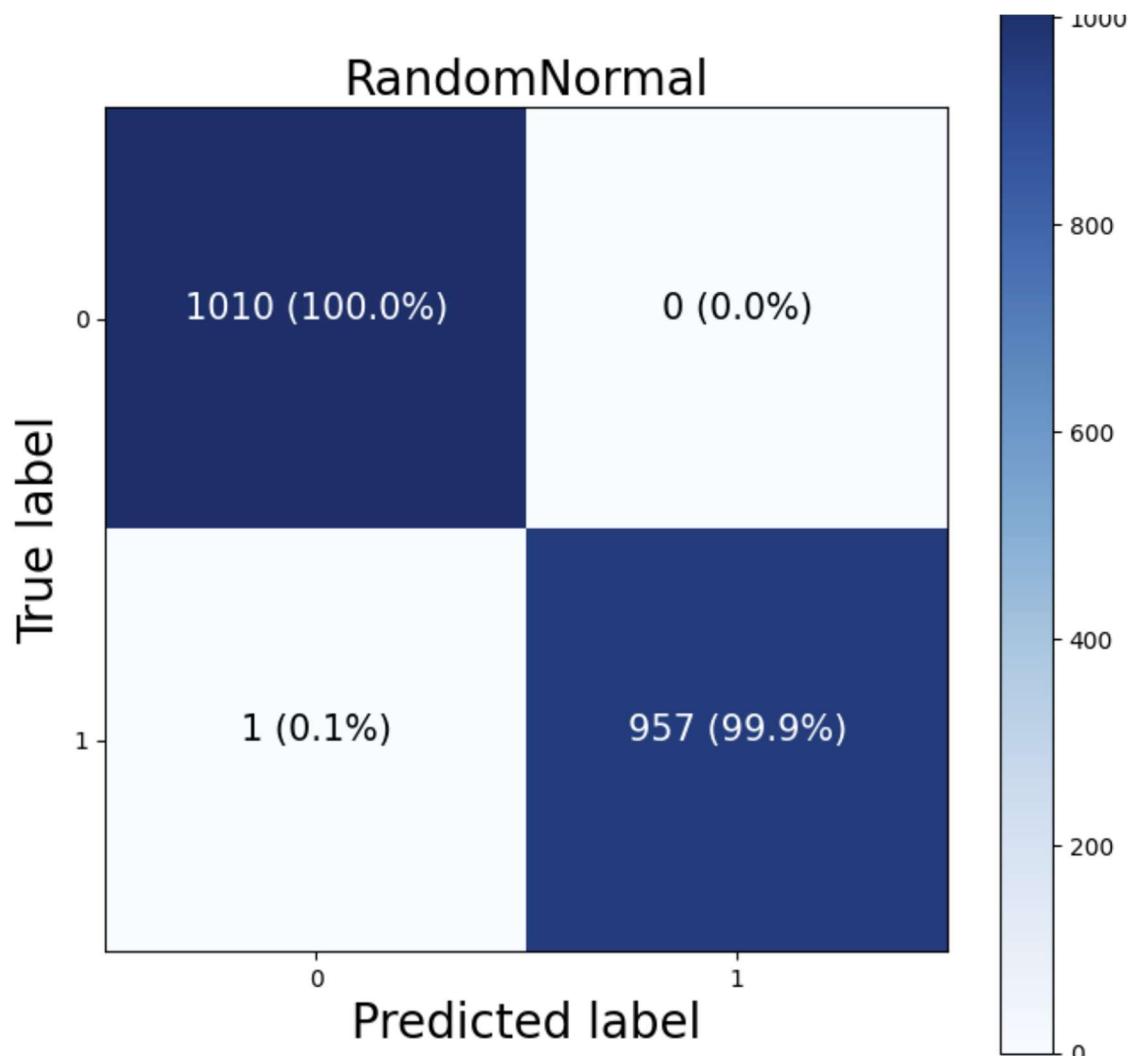
Model: "sequential\_4"

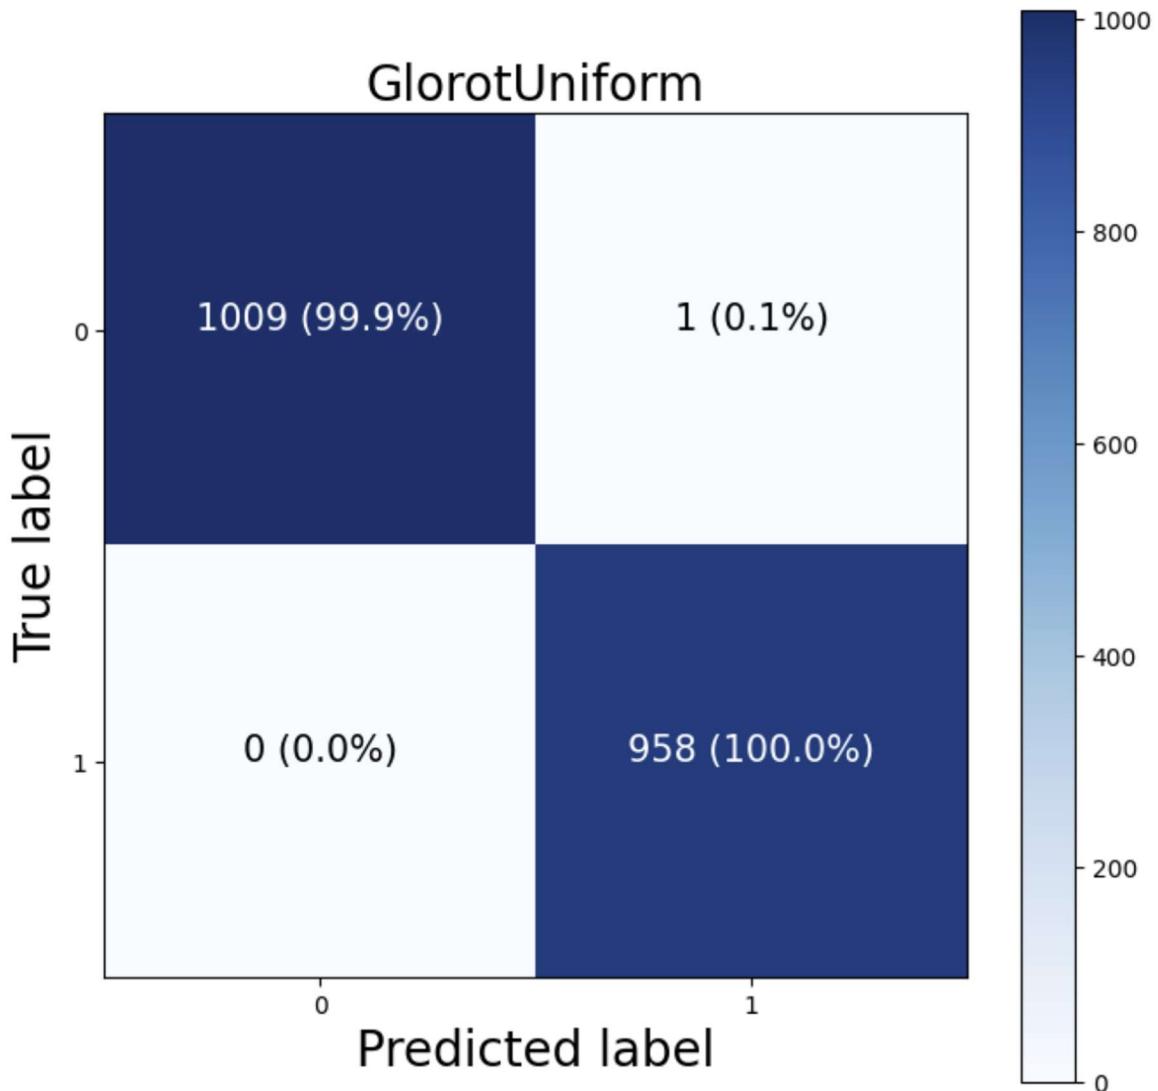
Layer (type)	Output Shape	Param #
<hr/>		
flatten_4 (Flatten)	(None, 784)	0
dense_8 (Dense)	(None, 128)	100480
dense_9 (Dense)	(None, 1)	129
<hr/>		
Total params: 100,609		
Trainable params: 100,609		
Non-trainable params: 0		

62/62 [=====] - 0s 869us/step

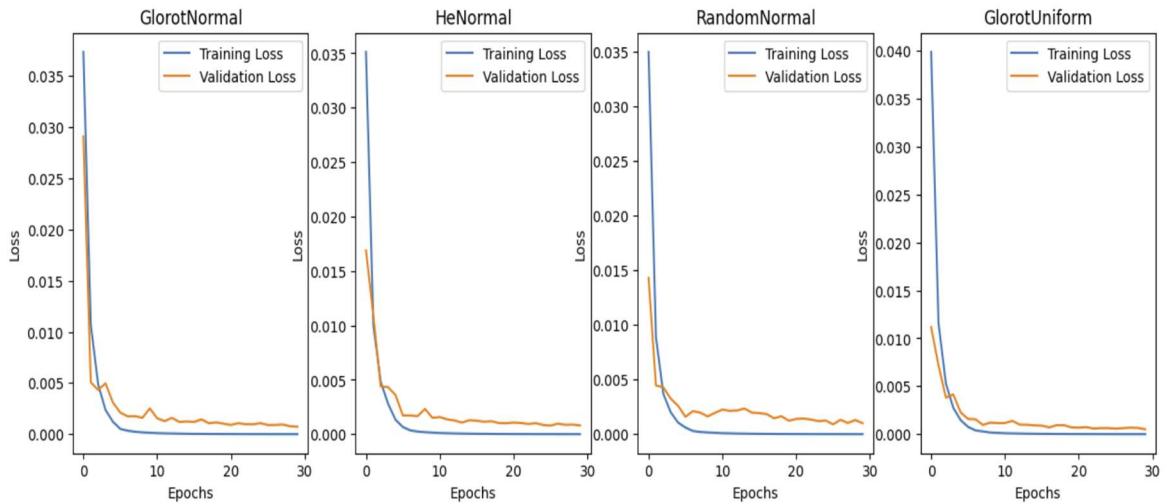






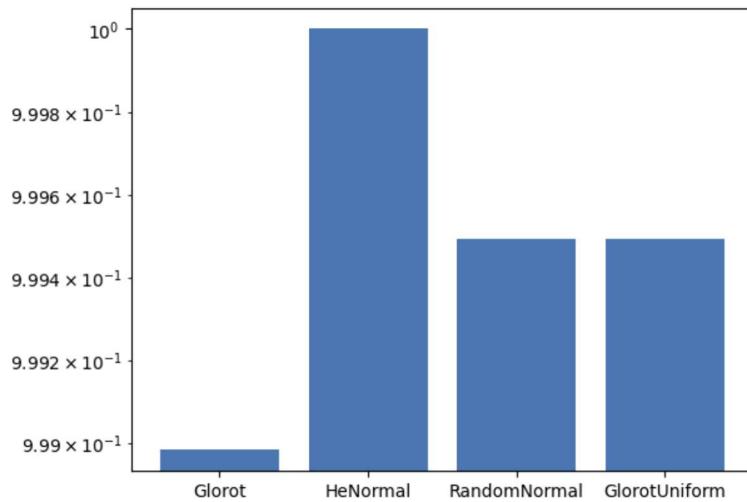


All of these initializers initializes the weights with a normal distribution centered around 0 but with difference in standard deviation. Glorot calculates standard deviation based on number of input units and output units, HE Normal calculates standard deviation with respect to number of input units and RandomNormal calculates std. as specified by the user. Here, He Normal performs the best with 100% accuracy. It performs better with ReLU activation function that has been used in the hidden layers.



**HeNormal shows a less erratic behavior as compared to Glorot and RandomNormal.**

**HeNormal is also smoother after 10 epochs as compared to other initializers. RandomNormal is the most erratic throughout the epoch cycle.**



In the above plot, HE Normal performs the best with the highest accuracy, followed by RandomNormal and GlorotUniform.

He Normal performs best with ReLU and Glorot performs best with sigmoid. Here, in the hidden layer ReLU is used, therefore, HeNormal performs better than the other.

**TABLE for 2-Class Hyperparameters – With Four Different Weight Initializers**

Activation function – hidden layer	ReLU
Activation function - output layer	Sigmoid
Weight Initializer	GlorotNormal(seed=1), HeNormal(seed=2), RandomNormal(mean=0.0, stddev=0.1, seed=3), GlorotUniform(seed=4)
Number of hidden layers	1 (2 including the output layer)
Neurons in hidden layers	128
Neurons in the output layer	1
Loss function	Binary_crossentropy
Optimizer	Adam
Number of epochs	20
Batch size	32
Learning rate	0.001
Evaluation Metric	Accuracy

## Q 3. 10 - class classification of MNIST

```
1 model.summary()
```

Model: "sequential\_6"

Layer (type)	Output Shape	Param #
<hr/>		
flatten_6 (Flatten)	(None, 784)	0
dense_12 (Dense)	(None, 128)	100480
dense_13 (Dense)	(None, 10)	1290
<hr/>		
Total params: 101,770		
Trainable params: 101,770		
Non-trainable params: 0		

```

]: 1 #iling the model .We use sparse categorical that will hot encode the classes internally.
2 compile(optimizer=tf.keras.optimizers.Adam(learning_rate=learning_rate), loss='sparse_categorical_crossentropy',metrics=['acc'])

]: 1 # Early stopping criteria based on validation loss
2 # patience = Number of epochs with no improvement after which training will be stopped.
3 # restore_best_weights = Whether to restore model weights from the epoch with the best value of the monitored quantity.
4 early_stopping = EarlyStopping(monitor='val_loss', patience=10, restore_best_weights=True)

]: 1 # Training the model
2 history = model.fit(x_train, y_train, batch_size=batch_size, epochs=20, validation_split=0.2, callbacks=[early_stopping])

Epoch 1/20
1500/1500 [=====] - 2s 1ms/step - loss: 0.2875 - accuracy: 0.9175 - val_loss: 0.1539 - val_accuracy: 0.9578
Epoch 2/20
1500/1500 [=====] - 2s 1ms/step - loss: 0.1271 - accuracy: 0.9634 - val_loss: 0.1077 - val_accuracy: 0.9690
Epoch 3/20
1500/1500 [=====] - 2s 1ms/step - loss: 0.0879 - accuracy: 0.9744 - val_loss: 0.0967 - val_accuracy: 0.9711

]: 1 # Evaluate the model on the test data
2 test_loss, test_accuracy = model.evaluate(x_test, y_test)
3 print("Test Accuracy:", test_accuracy)
4

313/313 [=====] - 0s 772us/step - loss: 0.0789 - accuracy: 0.9781
Test Accuracy: 0.9781000018119812

```

## CONFUSION MATRIX

```

[67]: 1 #draw_confusion_matrix(y_test, y_pred_classes,"classes")
2 # Construct a confusion matrix
3 conf_matrix = confusion_matrix(y_test, y_pred_classes)
4 print("Confusion Matrix:")
5 print(conf_matrix)
6

Confusion Matrix:
[[ 969   0   1   1   1   1   3   2   2   0]
 [  0 1120   4   2   0   1   3   1   4   0]
 [  2   0 1014   5   1   0   2   5   3   0]
 [  0   0   5 996   0   1   0   3   3   2]
 [  1   0   3   0 961   0   5   2   0  10]
 [  2   0   0  17   1 859   6   2   4   1]
 [  6   2   3   1   3   1 939   0   3   0]
 [  1   3  12   4   0   0   0 1002   3   3]
 [  3   0   5  11   4   1   1   4  942   3]
 [  1   3   0   6   7   3   0   9   1  979]]

```

As per the above confusion matrix, mostly all the classes from 0-9 are classified correctly. Class 0 and 1 are the least misclassified. Class 5 is misclassified as class 3 for 17 times; Class 7 has been misclaasified as 2 for 12 times; Class 8 is misclassified as 3 for 11 times.

```
1 # Display a classification report
2 class_report = classification_report(y_test, y_pred_classes)
3 print("Classification Report:")
4 print(class_report)
```

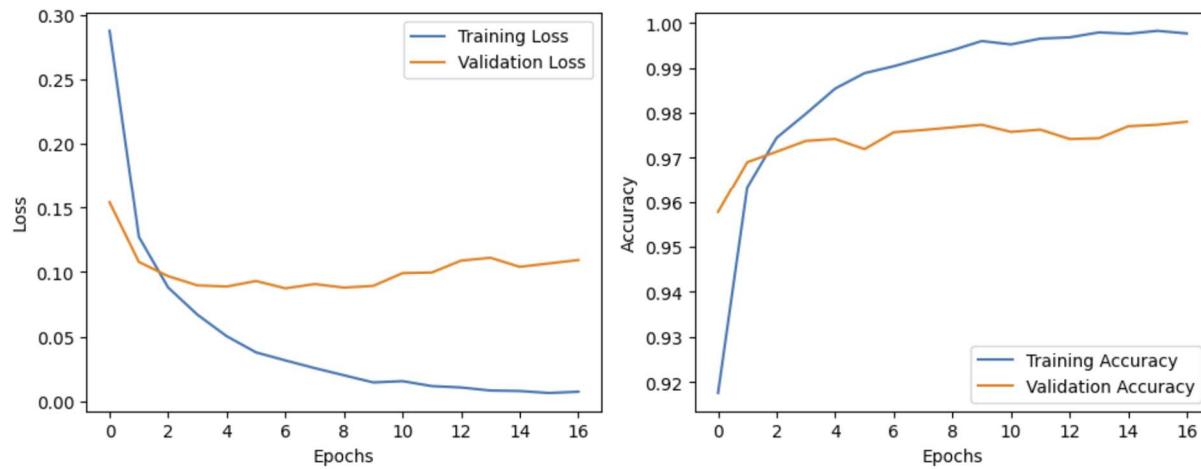
Classification Report:

	precision	recall	f1-score	support
0	0.98	0.99	0.99	980
1	0.99	0.99	0.99	1135
2	0.97	0.98	0.98	1032
3	0.95	0.99	0.97	1010
4	0.98	0.98	0.98	982
5	0.99	0.96	0.98	892
6	0.98	0.98	0.98	958
7	0.97	0.97	0.97	1028
8	0.98	0.97	0.97	974
9	0.98	0.97	0.98	1009
accuracy			0.98	10000
macro avg	0.98	0.98	0.98	10000
weighted avg	0.98	0.98	0.98	10000

```

1 # Plot the learning curve
2 plt.figure(figsize=(10, 4))
3 plt.subplot(1, 2, 1)
4 plt.plot(history.history['loss'], label='Training Loss')
5 plt.plot(history.history['val_loss'], label='Validation Loss')
6 plt.xlabel('Epochs')
7 plt.ylabel('Loss')
8 plt.legend()
9
10 plt.subplot(1, 2, 2)
11 plt.plot(history.history['accuracy'], label='Training Accuracy')
12 plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
13 plt.xlabel('Epochs')
14 plt.ylabel('Accuracy')
15 plt.legend()
16
17 plt.tight_layout()
18 plt.show()

```

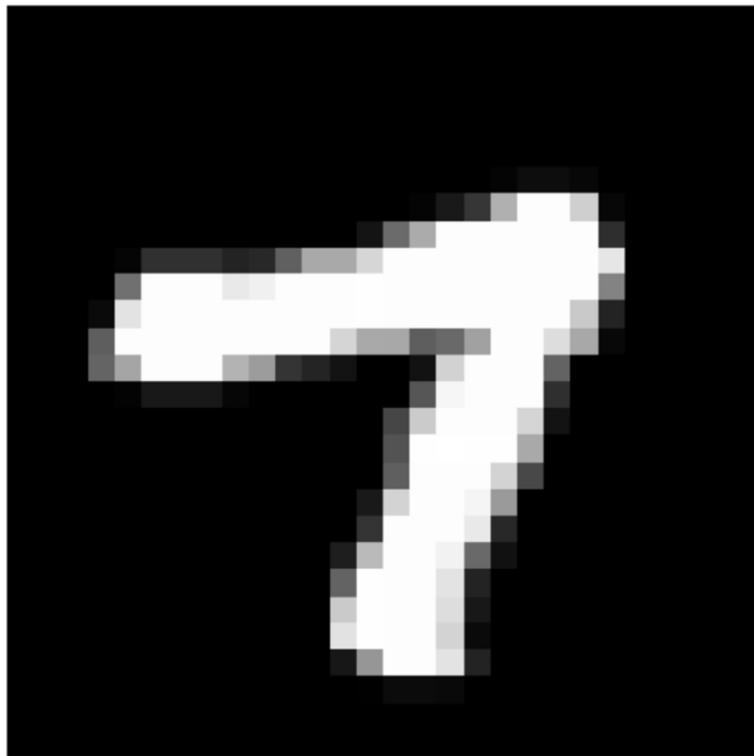


```

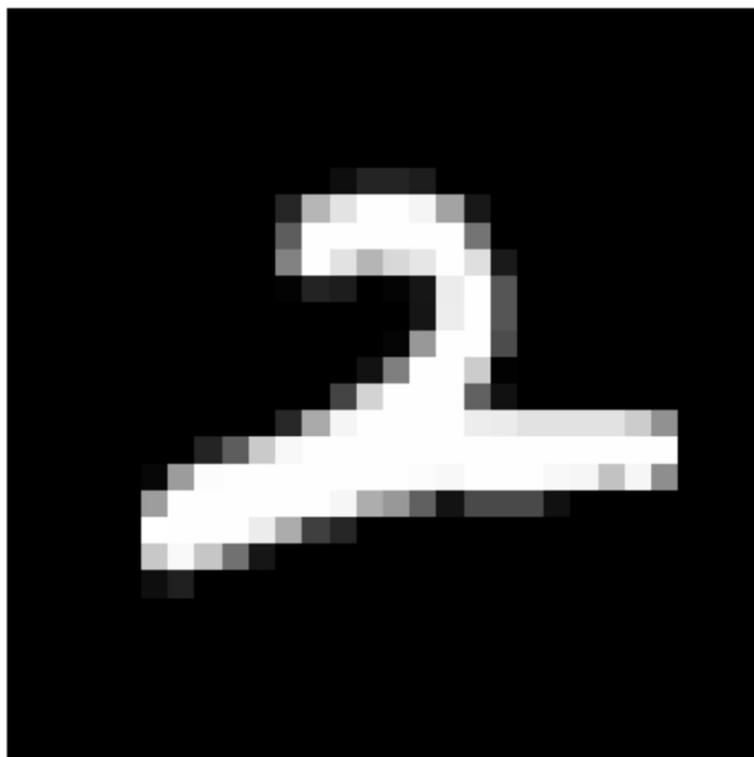
1 # Show some examples of predictions
2 eg_indices = np.random.choice(len(x_test), 10) # 10 randomly selected samples
3 eg_images = x_test[eg_indices]
4 eg_labels = y_test[eg_indices]
5 eg_predictions = model.predict(eg_images)
6 eg_predictions_classes = np.argmax(eg_predictions, axis=1)
7 for i in range(len(eg_images)):
8     plt.figure()
9     plt.imshow(eg_images[i], cmap='gray')
10    plt.title(f"True Label: {eg_labels[i]}, Predicted: {eg_predictions_classes[i]}")
11    plt.axis('off')
12

```

True Label: 7, Predicted: 7



True Label: 2, Predicted: 2



**TABLE for 10 Class hyperparameters**

Activation function – hidden layer	ReLU
Activation function - output layer	Softmax
Weight Initializer	Default - GlorotUniform
Number of hidden layers	1 (2 including the output layer)
Neurons in hidden layers	128
Neurons in the output layer	10
Loss function	sparse_categorical_crossentropy
Optimizer	Adam
Number of epochs	20
Batch size	32
Learning rate	0.001
Evaluation Metric	Accuracy